

Noam Simon 208388850

Chen Eliyahou 312490675

Report:

1. Summarize statistics regarding our experiments:
 - a. The various thresholds we employed:
 - i. We computed similarities only between words that occur at least 100 times in the corpus.
 - ii. We ignored words that have less than 75 occurrences.
 - iii. We limited the feature vector to 100 most common contexts per word.
 - iv. We ignored function word.